

Endnotes

1. Shafran, I., Cao, Y. et al., 2022, 'ReAct: Synergizing Reasoning and Acting in Language Models'. Available at: <https://arxiv.org/abs/2210.03629>.
2. Wei, J., Wang, X. et al., 2023, 'Chain-of-Thought Prompting Elicits Reasoning in Large Language Models'. Available at: <https://arxiv.org/pdf/2201.11903.pdf>.
3. Wang, X. et al., 2022, 'Self-Consistency Improves Chain of Thought Reasoning in Language Models'. Available at: <https://arxiv.org/abs/2203.11171>.
4. Diao, S. et al., 2023, 'Active Prompting with Chain-of-Thought for Large Language Models'. Available at: <https://arxiv.org/pdf/2302.12246.pdf>.
5. Zhang, H. et al., 2023, 'Multimodal Chain-of-Thought Reasoning in Language Models'. Available at: <https://arxiv.org/abs/2302.00923>.
6. Yao, S. et al., 2023, 'Tree of Thoughts: Deliberate Problem Solving with Large Language Models'. Available at: <https://arxiv.org/abs/2305.10601>.
7. Long, X., 2023, 'Large Language Model Guided Tree-of-Thought'. Available at: <https://arxiv.org/abs/2305.08291>.
8. Google. 'Google Gemini Application'. Available at: <http://gemini.google.com>.
9. Swagger. 'OpenAPI Specification'. Available at: <https://swagger.io/specification/>.
10. Xie, M., 2022, 'How does in-context learning work? A framework for understanding the differences from traditional supervised learning'. Available at: <https://ai.stanford.edu/blog/understanding-incontext/>.
11. Google Research. 'ScaNN (Scalable Nearest Neighbors)'. Available at: <https://github.com/google-research/google-research/tree/master/scann>.
12. LangChain. 'LangChain'. Available at: <https://python.langchain.com/v0.2/docs/introduction/>.
13. Sokratis Kartakis, 2024, 'GenAI in Production: MLOps or GenAIOps?'. Available at: <https://medium.com/google-cloud/genai-in-production-mlops-or-genaiops-25691c9becd0>.
14. Sokratis Kartakis, 2024 'Gen AI Ops, Operationalize Generative AI, A practical Guide'. Available at: <https://medium.com/google-cloud/genaiops-operationalize-generative-ai-a-practical-guide-d5bedaa59d78>.

15. Cloud Trace overview. Available at: <https://cloud.google.com/trace/docs/overview>.
16. Berkeley Function-Calling Leaderboard (BFCL). Available at: https://gorilla.cs.berkeley.edu/blogs/8_berkeley_function_calling_leaderboard.html.
17. Karthik Narasimhan, et al. 2024, 'τ-bench'. Available at <https://arxiv.org/abs/2406.12045>.
18. Karthik Valmeekam, et al., 2023, 'PlanBench'. Available at: <https://arxiv.org/abs/2206.10498>.
19. Xiao Liu, et al., 2023, 'AgentBench'. Available at: <https://arxiv.org/abs/2308.03688>.
20. Martin Iglesias, et al., 2025, 'DBASStep' Available at: <https://huggingface.co/spaces/adyen/DABstep>.
21. LangSmith platform for agent observability.
Available at: <https://docs.smith.langchain.com/evaluation/concepts#agents>.
22. Mingchen Zhuge, et al., 2024, 'Agent-as-a-Judge: Evaluate Agents with Agents'.
Available at: <https://arxiv.org/abs/2410.10934>.
23. Multi-agent documentation from LangGraph.
Available at: https://langchain-ai.github.io/langgraph/concepts/multi_agent/.
24. LangChain blog 2024, 'Multi-agent workflows'.
Available at: <https://blog.langchain.dev/langgraph-multi-agent-workflows/>.
25. Vectorize blog 2024, 'How I finally got agentic RAG to work right'.
Available at: <https://vectorize.io/how-i-finally-got-agentic-rag-to-work-right/>.
26. Vertex AI Search, product documentation. Available at: <https://cloud.google.com/enterprise-search>.
27. Vertex AI Search Builder APIs, product documentation.
Available at: <https://cloud.google.com/generative-ai-app-builder/docs/builder-apis>.
28. Vertex AI RAG Engine, product documentation.
Available at: <https://cloud.google.com/vertex-ai/generative-ai/docs/rag-overview>.
29. Agentspace product documentation.
Available at: <https://cloud.google.com/agentspace/agentspace-enterprise/docs/overview>.
30. NotebookLM Enterprise product documentation.
Available at: <https://cloud.google.com/agentspace/notebooklm-enterprise/docs/overview>.

31. Juraj Gottweis, et. al., 2025, 'Accelerating scientific breakthroughs with an AI co-scientist'. Available at: <https://research.google/blog/accelerating-scientific-breakthroughs-with-an-ai-co-scientist/>.
32. Hamsa Buvaraghan, et al. 2025, 'Announcing public beta of Gen AI Toolbox for Databases'. Available at: <https://cloud.google.com/blog/products/ai-machine-learning/announcing-gen-ai-toolbox-for-databases-get-started-today?e=48754805>.
33. Google Cloud Integration Connectors, product documentation.
Available at: <https://cloud.google.com/integration-connectors/docs>.
34. Apigee API Hub, product documentation.
Available at: <https://cloud.google.com/apigee/docs/apihub/what-is-api-hub>.
35. Vertex AI Model Garden, product documentation.
Available at: <https://cloud.google.com/model-garden>.
36. Gemini family of LLMs, product documentation.
Available at: <https://cloud.google.com/vertex-ai/generative-ai/docs/learn/models#gemini-models>.
37. Get Started Evaluating Agents with the Vertex Eval Service. Available at: <https://cloud.google.com/vertex-ai/generative-ai/docs/models/evaluation-agents>.
38. Irina Sigler, Ivan Nardini. Jan 2025 'Introducing Agent Evaluation in Vertex AI'. Available at: <https://cloud.google.com/blog/products/ai-machine-learning/introducing-agent-evaluation-in-vertex-ai-gen-ai-evaluation-service?e=48754805>.
39. Review sample agent evaluation notebooks for LangGraph, CrewAI, and LangChain.
Available at: <https://github.com/GoogleCloudPlatform/generative-ai/blob/main/gemini/evaluation/>.
40. Review many sample agents, primarily beginner and intermediate level.
Available at: <https://github.com/GoogleCloudPlatform/generative-ai/>.
41. Review many sample agents, intermediate and advanced levels.
Available at: <https://github.com/GoogleCloudPlatform/applied-ai-engineering-samples>.